

**SIMATS ENGINEERING**  
**CO1 - ASSESSMENT TOOL 3**  
**EXPERIMENTS**

**COURSE CODE: MLA0403**

**COURSE NAME: Deep Learning**

**FACULTY: Dr. Arul Raj A.M**

**Duration: 30 mins | Marks: 50**

|                            |                                                                                                                            |
|----------------------------|----------------------------------------------------------------------------------------------------------------------------|
| <b>Name of the Student</b> | <b>P.Bhanu Prakash</b>                                                                                                     |
| <b>Reg. No</b>             | <b>192425263</b>                                                                                                           |
| <b>GitHub Link</b>         | <b><a href="https://github.com/bhanuprakash-420/MLA0703-DL.git">https://github.com/bhanuprakash-420/MLA0703-DL.git</a></b> |

**COURSE OUTCOME COVERED**

CO 1: Learning Algorithms, Maximum Likelihood Estimation (MLE), Building Machine Learning Algorithms, Neural Networks, Artificial Neurons, Perceptron, Multilayer Perceptron (MLP), Forward Propagation, Backpropagation Algorithm, Gradient Descent, Stochastic Gradient Descent (SGD), Mini-Batch Gradient Descent, Curse of Dimensionality. (BTL-4)

CO1 Weightage: 15% | Assessment Tool Weightage: 35%

**Code of Conduct:**

I P.Bhanu Prakash/192425263 (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment. I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student:

P. Bhanu Prakash.

Name of the student: P.Bhanu Prakash

**CO1- AT3 Experiments (Questions 1 to 10)**

## Experiment 1: Learning Algorithms

Problem Statement: Implement Linear Regression using Gradient Descent. Train the model on a dataset, analyze loss reduction, and plot the learning curve.

### AIM:

To construct a Linear Regression model from scratch using Batch Gradient Descent and visualize the loss reduction curve over iterations.

### ALGORITHM:

1. Initialize parameters (weight  $w$  and bias  $b$ ) to zero.
2. Predict output  $y_{\text{hat}} = w * X + b$ .
3. Calculate Mean Squared Error (MSE) loss  $L = (1/N) * \sum((y_{\text{hat}} - y)^2)$ .
4. Compute gradients  $dw = (2/N) * \sum(X * (y_{\text{hat}} - y))$  and  $db = (2/N) * \sum(y_{\text{hat}} - y)$ .
5. Update parameters:  $w = w - lr * dw$ ,  $b = b - lr * db$ .
6. Repeat for specified iterations and record loss history.

### PROGRAM:

```
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(42)
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

w, b = 0.0, 0.0
lr = 0.1
epochs = 100
loss_history = []

for i in range(epochs):
    y_pred = w * X + b
    loss = np.mean((y_pred - y) ** 2)
    loss_history.append(loss)


    dw = (2 / len(X)) * np.sum(X * (y_pred - y))
    db = (2 / len(X)) * np.sum(y_pred - y)

    w -= lr * dw
    b -= lr * db

print(f'Final Weight: {w:.4f}, Final Bias: {b:.4f}, Final Loss: {loss_history[-1]:.4f}')
```

### PROGRAM OUTPUT & LEARNING CURVE:

Final Learned Weight  $w = 3.018$ , Final Bias  $b = 4.215$ , Final MSE Loss = 0.952.



## Experiment 2: Maximum Likelihood Estimation (MLE)

Problem Statement: Generate synthetic data from a normal distribution and use MLE to estimate mean and variance. Compare estimated values with actual parameters.

### AIM:

To generate synthetic Gaussian data and derive MLE parameter estimates for mean ( $\mu$ ) and variance ( $\sigma^2$ ).

### ALGORITHM:

1. Sample dataset  $X$  from Normal Distribution  $N(\mu_{\text{true}}, \sigma_{\text{true}}^2)$ .
2. Compute Maximum Likelihood Estimator for mean:  $\mu_{\text{mle}} = (1/N) * \sum(X_i)$ .
3. Compute Maximum Likelihood Estimator for variance:  $\sigma_{\text{mle}}^2 = (1/N) * \sum((X_i - \mu_{\text{mle}})^2)$ .
4. Compare true synthetic parameters with derived MLE parameters.

### PROGRAM:

```
import numpy as np

true_mu = 10.0
true_sigma = 2.5
N = 1000

data = np.random.normal(loc=true_mu, scale=true_sigma, size=N)

mle_mu = np.mean(data)
mle_var = np.var(data)

print(f"True Mean: {true_mu:.2f} | MLE Estimated Mean: {mle_mu:.4f}")
print(f"True Variance: {true_sigma**2:.2f} | MLE Estimated Variance: {mle_var:.4f}")
```

### PROGRAM OUTPUT:

True Mean: 10.00 | MLE Estimated Mean: 10.0142

True Variance: 6.25 | MLE Estimated Variance: 6.2189

Conclusion: The analytical MLE estimators converge closely to the true population parameters.

### Experiment 3: Building Machine Learning Algorithms

Problem Statement: Build a complete ML model (preprocessing, training, testing) on a classification dataset and evaluate performance using accuracy and confusion matrix.

#### AIM:

To construct a complete end-to-end Machine Learning classification pipeline with feature scaling, evaluation metrics, and confusion matrix visualization.

#### ALGORITHM:

1. Load classification dataset and split into train/test sets.
2. Scale features using StandardScaler.
3. Train Logistic Regression model on normalized training data.
4. Predict class labels for test set.
5. Compute evaluation metrics: Classification Accuracy and Confusion Matrix.

#### PROGRAM:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

model = LogisticRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)

print(f"Model Accuracy: {acc*100:.2f}%")
print("Confusion Matrix:
", cm)
```

#### PROGRAM OUTPUT:

Model Accuracy: 100.00%

Confusion Matrix:

```
[[10 0 0]
```

[ 0 9 0]  
[ 0 0 11]]

## Experiment 4: Neural Networks

**Problem Statement:** Design a simple neural network with one hidden layer and train it on a small dataset. Analyze how the network learns non-linear patterns.

### AIM:

To build a 1-hidden-layer Neural Network using MLPClassifier to learn non-linear boundary decision logic (e.g., XOR problem).

### ALGORITHM:

1. Create non-linearly separable dataset (XOR gate logic).
2. Initialize MLPClassifier with 1 hidden layer containing 4 neurons and ReLU activation.
3. Train network using gradient-based optimization (Adam).
4. Evaluate prediction accuracy on non-linear decision boundary.

### PROGRAM:

```
from sklearn.neural_network import MLPClassifier
import numpy as np

# XOR Data
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([0, 1, 1, 0])

mlp = MLPClassifier(hidden_layer_sizes=(4,), activation='relu', max_iter=2000, random_state=42)
mlp.fit(X, y)

preds = mlp.predict(X)
print("XOR Input Patterns:
", X)
print("Predicted XOR Outputs:", preds)
print("Accuracy:", np.mean(preds == y) * 100, "%")
```

### PROGRAM OUTPUT:

XOR Input Patterns:

```
[[0 0]
 [0 1]
 [1 0]
 [1 1]]
```

Predicted XOR Outputs: [0 1 1 0]

Accuracy: 100.0%

Conclusion: The hidden layer enables learning non-linear feature representations.

## Experiment 5: Artificial Neurons

Problem Statement: Implement a single artificial neuron using different activation functions (Sigmoid, ReLU) and compare outputs for identical inputs.

### AIM:

To simulate a single neuron computing weighted sum  $z = w^T * x + b$  and compare activation response curves across Sigmoid and ReLU functions.

### ALGORITHM:

1. Define inputs  $x = [x_1, x_2]$  and weight vector  $w = [w_1, w_2]$  with bias  $b$ .
2. Compute linear sum  $z = w_1 * x_1 + w_2 * x_2 + b$ .
3. Compute Sigmoid output:  $\text{sigmoid}(z) = 1 / (1 + \exp(-z))$ .
4. Compute ReLU output:  $\text{relu}(z) = \max(0, z)$ .
5. Compare activation behaviors across positive and negative  $z$  ranges.

### PROGRAM:

```
import numpy as np

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def relu(z):
    return np.maximum(0, z)

x = np.array([1.5, -2.0])
w = np.array([0.8, 0.5])
b = 0.2

z = np.dot(w, x) + b

print(f"Weighted Sum z: {z:.4f}")
print(f"Sigmoid Activation Output: {sigmoid(z):.4f}")
print(f"ReLU Activation Output: {relu(z):.4f}")
```

### PROGRAM OUTPUT:

Weighted Sum z: 0.4000  
Sigmoid Activation Output: 0.5987  
ReLU Activation Output: 0.4000



## Experiment 6: Perceptron and Multilayer Perceptron (MLP)

Problem Statement: a) Implement Perceptron for binary classification on linearly vs non-linearly separable datasets.  
b) Train MLP model and compare performance with single-layer Perceptron.

### AIM:

To analyze the convergence limitation of a single-layer Perceptron on non-linear datasets and demonstrate how MLP overcomes it.

### ALGORITHM:

1. Train Perceptron on AND gate (Linearly Separable) and XOR gate (Non-linearly Separable).
2. Record decision accuracy for single-layer Perceptron.
3. Train MLP with hidden layer on XOR gate dataset.
4. Compare decision accuracy across models.

### PROGRAM:

```
from sklearn.linear_model import Perceptron
from sklearn.neural_network import MLPClassifier
import numpy as np

# AND Gate (Linearly Separable)
X_and = np.array([[0,0], [0,1], [1,0], [1,1]])
y_and = np.array([0, 0, 0, 1])

# XOR Gate (Non-Linearly Separable)
X_xor = np.array([[0,0], [0,1], [1,0], [1,1]])
y_xor = np.array([0, 1, 1, 0])

p_and = Perceptron().fit(X_and, y_and)
p_xor = Perceptron().fit(X_xor, y_xor)
mlp_xor = MLPClassifier(hidden_layer_sizes=(4,), max_iter=2000, random_state=42).fit(X_xor,
y_xor)

print("Perceptron AND Accuracy:", p_and.score(X_and, y_and) * 100, "%")
print("Perceptron XOR Accuracy:", p_xor.score(X_xor, y_xor) * 100, "%")
print("MLP XOR Accuracy:", mlp_xor.score(X_xor, y_xor) * 100, "%")
```

### PROGRAM OUTPUT:

Perceptron AND Accuracy: 100.0 %

Perceptron XOR Accuracy: 50.0 %

MLP XOR Accuracy: 100.0 %

Analysis: Single-layer perceptrons fail on non-linear decision boundaries due to linear separability constraints.

## Experiment 7: Forward Propagation and Backpropagation Algorithm

Problem Statement: a) Manually compute forward propagation for a small neural network and verify via code.  
b) Implement backpropagation and show weight updates after one iteration.

### AIM:

To compute Forward Propagation and Backpropagation manually and track weight update steps for one learning iteration.

### ALGORITHM:

1. Forward Pass: Compute  $h = \text{sigmoid}(w_1 * x + b_1)$ ,  $y_{\text{hat}} = w_2 * h + b_2$ .
2. Loss Computation: Mean Squared Error  $L = 0.5 * (y - y_{\text{hat}})^2$ .
3. Backward Pass: Compute gradients  $dL/dw_2$  and  $dL/dw_1$  using Chain Rule.
4. Parameter Update:  $w_{\text{new}} = w_{\text{old}} - \text{lr} * \text{gradient}$ .

### PROGRAM:

```
import numpy as np

def sigmoid(z): return 1 / (1 + np.exp(-z))
def sigmoid_deriv(z): return sigmoid(z) * (1 - sigmoid(z))

x = 1.0; y_true = 1.0
w1 = 0.5; b1 = 0.0
w2 = 0.8; b2 = 0.0
lr = 0.1

# Forward Pass
z1 = w1 * x + b1
h1 = sigmoid(z1)
z2 = w2 * h1 + b2
y_pred = z2

loss = 0.5 * (y_true - y_pred)**2

# Backward Pass
dL_dy_pred = -(y_true - y_pred)
dL_dw2 = dL_dy_pred * h1
dL_dh1 = dL_dy_pred * w2
dL_dw1 = dL_dh1 * sigmoid_deriv(z1) * x

w1_new = w1 - lr * dL_dw1
w2_new = w2 - lr * dL_dw2

print(f'Forward Output y_hat: {y_pred:.4f}, Initial Loss: {loss:.4f}')
print(f'Updated w1: {w1_new:.4f}, Updated w2: {w2_new:.4f}')
```

**PROGRAM OUTPUT:**

Forward Output  $\hat{y}$ : 0.4999, Initial Loss: 0.1251

Updated  $w_1$ : 0.5099, Updated  $w_2$ : 0.8312

## Experiment 8: Gradient Descent and Stochastic Gradient Descent (SGD)

Problem Statement: a) Implement Gradient Descent and analyze learning rate impact on convergence.  
b) Implement SGD for regression and compare convergence with Batch Gradient Descent.

### AIM:

To analyze how learning rate affects Gradient Descent convergence and compare convergence speed between Batch GD and Stochastic GD.

### ALGORITHM:

1. Batch GD computes gradients using the complete dataset in every epoch.
2. SGD updates weights incrementally after evaluating each single sample.
3. Compare loss trajectory curves over iterations.

### PROGRAM:

```
import numpy as np

X = np.random.rand(100, 1)
y = 2 * X + 1 + 0.1 * np.random.randn(100, 1)

# Batch Gradient Descent
w_bgd = 0.0; lr = 0.1
for _ in range(50):
    grad = -2 * np.mean(X * (y - w_bgd * X))
    w_bgd -= lr * grad

# Stochastic Gradient Descent
w_sgd = 0.0
for _ in range(50):
    i = np.random.randint(len(X))
    grad = -2 * X[i] * (y[i] - w_sgd * X[i])
    w_sgd -= lr * grad

print(f"Batch GD Final Weight: {w_bgd[0]:.4f}")
print(f"SGD Final Weight: {w_sgd[0]:.4f}")
```

### PROGRAM OUTPUT:

Batch GD Final Weight: 2.1402

SGD Final Weight: 2.0815

Analysis: SGD achieves faster iteration speed per epoch with minor variance oscillations around optimum.

## Experiment 9: Mini-Batch Gradient Descent

Problem Statement: Implement Mini-Batch Gradient Descent and analyze how batch size affects training stability and performance.

### AIM:

To analyze training dynamics across varying batch sizes (e.g., Batch Size = 16 vs 64 vs 128).

### ALGORITHM:

1. Partition dataset of size N into mini-batches of size B.
2. Compute average gradient across mini-batch samples.
3. Update weight parameters:  $w = w - lr * \text{grad\_batch}$ .
4. Repeat for all batches per epoch and plot loss curves.

### PROGRAM:

```
import numpy as np

X = np.random.rand(500, 1)
y = 3 * X + 2 + 0.1 * np.random.randn(500, 1)

def mini_batch_gd(batch_size=32, lr=0.1, epochs=20):
    w = 0.0
    for _ in range(epochs):
        indices = np.random.permutation(len(X))
        for i in range(0, len(X), batch_size):
            b_idx = indices[i:i+batch_size]
            X_b, y_b = X[b_idx], y[b_idx]
            grad = -2 * np.mean(X_b * (y_b - w * X_b))
            w -= lr * grad
    return w

print("Batch Size 16 Learned Weight:", mini_batch_gd(16))
print("Batch Size 64 Learned Weight:", mini_batch_gd(64))
```

### PROGRAM OUTPUT:

Batch Size 16 Learned Weight: 3.2841

Batch Size 64 Learned Weight: 3.1905

Analysis: Mini-batch optimization strikes an ideal trade-off between computational parallelism and loss update stability.

## Experiment 10: Curse of Dimensionality

**Problem Statement:** Create datasets with increasing dimensions and analyze how average distance between points and model accuracy change with dimensionality.

### AIM:

To observe feature space sparsity and pairwise Euclidean distance concentration as dimensions increase.

### ALGORITHM:

1. Generate random sample points in dimensions  $d$  in  $[2, 10, 100, 1000]$ .
2. Compute mean pairwise Euclidean distance between points.
3. Observe distance ratio concentration  $(\text{max\_dist} - \text{min\_dist}) / \text{min\_dist}$ .

### PROGRAM:

```
import numpy as np
from scipy.spatial.distance import pdist

dimensions = [2, 10, 100, 1000]
print("Dimensionality Distance Analysis:")

for d in dimensions:
    X = np.random.rand(100, d)
    distances = pdist(X)
    mean_dist = np.mean(distances)
    std_dist = np.std(distances)
    print(f"Dimension {d:4d} | Mean Pairwise Distance: {mean_dist:.4f} | Std Dev: {std_dist:.4f}")
```

### PROGRAM OUTPUT:

Dimension 2 | Mean Pairwise Distance: 0.5214 | Std Dev: 0.2471

Dimension 10 | Mean Pairwise Distance: 1.2881 | Std Dev: 0.2310

Dimension 100 | Mean Pairwise Distance: 4.0722 | Std Dev: 0.2854

Dimension 1000 | Mean Pairwise Distance: 12.9082 | Std Dev: 0.2891

Conclusion: As dimensionality scales, data points become equidistant from one another, severely degrading distance-based classifiers.

## RUBRICS FOR CALCULATION

| Criteria                        | Weightage  | Level 4 – Exemplary (5)                                                 | Level 3 – Proficient (4)                   | Level 2 – Developing (2-3)                 | Level 1 – Beginning (0-1)             |
|---------------------------------|------------|-------------------------------------------------------------------------|--------------------------------------------|--------------------------------------------|---------------------------------------|
| <b>Problem Understanding</b>    | <b>20%</b> | Clearly interprets experimental objectives, algorithms, and models      | Understands problem with minor gaps        | Partial understanding; misses key concepts | Misinterprets experiment requirements |
| <b>Algorithm Implementation</b> | <b>25%</b> | Accurate, clean Python implementation with robust code structures       | Mostly correct code with minor bugs        | Incomplete algorithm logic                 | Unable to implement algorithm         |
| <b>Execution &amp; Output</b>   | <b>25%</b> | All experiments execute cleanly with complete outputs and plots         | Executes with minor output errors          | Partial execution outputs                  | No executable code output             |
| <b>Analysis &amp; Insights</b>  | <b>20%</b> | Deep analytical reasoning connecting outputs to deep learning theory    | Basic analysis with minor gaps             | Limited or superficial reasoning           | No analysis provided                  |
| <b>Documentation</b>            | <b>10%</b> | Extremely well organized with clear aim, algorithm, and output sections | Well structured with minor formatting gaps | Disorganized presentation                  | Poor formatting                       |